

Problem 1: Logistic Regression

1.1:

In [9]:


```
import numpy as np
import mltools as ml
import matplotlib.pyplot as plt

iris = np.genfromtxt("iris.txt", delimiter=None)
X, Y = iris[:, 0:2], iris[:, -1] # get first two features & target
X, Y = ml.shuffleData(X, Y)      # reorder randomly rather than by class label
X, _ = ml.transforms.rescale(X)  # rescale to improve numerical stability, speed convergence


XA, YA = X[Y<2, :], Y[Y<2]     # Dataset A: class 0 vs class 1
XB, YB = X[Y>0, :], Y[Y>0]     # Dataset B: class 1 vs class 2

print("Dataset A:")
ml.plotClassify2D(None, XA, YA)
plt.show()
print("Dataset B:")
ml.plotClassify2D(None, XB, YB)
```

Dataset A:



Dataset B:



The first dataset is much more linearly separable because there is a gap between the different colored points.

1.2:


In [11]:

```
import mltools as ml
from logisticClassify2 import *
```


```
def myPlotBoundary(self, X,Y):
    """ Plot the (linear) decision boundary of the classifier, along with data """
    if len(self.theta) != 3: raise ValueError('Data & model must be 2D');
    ax = X.min(0),X.max(0); ax = (ax[0][0],ax[1][0],ax[0][1],ax[1][1]);
    ## TODO: find points on decision boundary defined by theta0 + theta1 X1 + theta2 X2 == 0
    x1b = np.array([ax[0],ax[1]]); # at X1 = points in x1b
    x2b = np.array([( -self.theta[1]*x1b[0]-self.theta[0])/self.theta[2],(-self.theta[1]*x1b[1]-self
.theta[0])/self.theta[2]]) # TODO find x2 values as a function of x1's values
    ## Now plot the data and the resulting boundary:
    A = Y==self.classes[0]; # and plot it:
    plt.plot(X[A,0],X[A,1], 'b.',X[~A,0],X[~A,1], 'r.',x1b,x2b,'k-'); plt.axis(ax); plt.draw();

print("Decision Boundary of Dataset A:")
learner = logisticClassify2()
learner.classes = np.unique(YA) # store the class values for this problem
wts = np.array([2,6,-1]); # TODO: fill in values
learner.theta = wts;
myPlotBoundary(learner,XA,YA)
plt.show()
print("Decision Boundary of Dataset B:")
learners = logisticClassify2()
learners.classes = np.unique(YB) # store the class values for this problem
wts = np.array([2,6,-1]); # TODO: fill in values
learners.theta = wts;
myPlotBoundary(learners,XB,YB)
```

Decision Boundary of Dataset A:



Decision Boundary of Dataset B:



1.3:

In [13]:

```
# Should go in your logistic2 class:
def myPredict(self,X):
    """ Return the predicted class of each data point in X"""
    Yhat = []
    a = 0
    for v in X:
```

```

    o = self.theta[2]*X[a,1] + self.theta[1]*X[a,0] + self.theta[0]
    if o > 0:
        Yhat.append(self.classes[1])
    else:
        Yhat.append(self.classes[0])
    a += 1
    ## TODO: compute linear response r[i] = theta0 + theta1 X[i,1] + theta2 X[i,2] for each i
    ## TODO: if r[i] > 0, predict class 1: Yhat[i] = self.classes[1]
    ##         else predict class 0: Yhat[i] = self.classes[0]
    return np.array(Yhat)

# Update our shell classifier definition
class logisticClassify2(ml.classifier):
    classes = []
    theta = np.array([-1, 0, 0]) # initialize theta to something
    plotBoundary = myPlotBoundary #
    predict = myPredict
    train = None

learnerA = logisticClassify2()
learnerA.classes = np.unique(YA) # store the class values for this problem
learnerA.theta = np.array([2,6,-1]); # TODO: insert hard-coded values
print("Dataset A: ")
print(learnerA.err(XA, YA))
learnerB = logisticClassify2()
learnerB.classes = np.unique(YB) # store the class values for this problem
learnerB.theta = np.array([2,6,-1]);
print("Dataset B: ")
print(learnerB.err(XB, YB))

```

Dataset A:
0.06060606060606061
Dataset B:
0.45454545454545453



1.4:


In [14]:

```

ml.plotClassify2D(learnerA, XA, YA)
plt.show()
ml.plotClassify2D(learnerB, XB, YB)
plt.show()

```





1.5:

$$J_j(\theta) = -y^{(j)} \log(\sigma(x^{(j)} \cdot \theta)) - (1 - y^{(j)}) \log(1 - \sigma(x^{(j)} \cdot \theta)) = -y^{(j)} \left(\frac{1}{\sigma(x^{(j)} \cdot \theta)} \sigma(x^{(j)} \cdot \theta) (1 - \sigma(x^{(j)} \cdot \theta)) \right) x^{(j)} - (1 - y^{(j)}) \left(\frac{1}{1 - \sigma(x^{(j)} \cdot \theta)} (1 - \sigma(x^{(j)} \cdot \theta)) \right) x^{(j)}$$

Cancel out similar variables = $-y^{(j)} x^{(j)} (1 - \sigma(x^{(j)} \cdot \theta)) - x^{(j)} (1 - y^{(j)}) - \sigma(x^{(j)} \cdot \theta) = x^{(j)}(-y^{(j)} + \sigma(x^{(j)} \cdot \theta)) + \sigma(x^{(j)} \cdot \theta) - \sigma(x^{(j)} \cdot \theta) = x^{(j)}(\sigma(x^{(j)} \cdot \theta) - y^{(j)})$

1.6:

In [25]:

```
def sigmoid(x):
    return 1. / (1 + np.exp(-x))

def myTrain(self, X, Y, initStep=1.0, stopTol=1e-4, stopEpochs=5000, plot=None):
    """ Train the logistic regression using stochastic gradient descent """
    from IPython import display
    M, N = X.shape; # initialize the model if necessary:
    self.classes = np.unique(Y); # Y may have two classes, any values
    XX = np.hstack((np.ones((M,1)), X)) # XX is X, but with an extra column of ones
    YY = ml.toIndex(Y, self.classes); # YY is Y, but with canonical values 0 or 1
    if len(self.theta) != N+1: self.theta = np.random.rand(N+1);
    # init loop variables:
    epoch=0; done=False; Jnll=[]; J01=[];
    while not done:
        stepsize, epoch = initStep*2.0/(2.0+epoch), epoch+1; # update stepsize
        # Do an SGD pass through the entire data set:
        for i in np.random.permutation(M):
            ri = self.theta[2] * X[i, 1] + self.theta[1] * X[i, 0] + self.theta[0] # TODO: compute linear response r(x)
            gradilist = [ sigmoid(ri) - YY[i], X[i,0] * (sigmoid(ri) - YY[i]), X[i,1] * (sigmoid(ri) - YY[i])]
            gradi = np.array(gradilist) # TODO: compute gradient of NLL loss
            self.theta -= stepsize * gradi; # take a gradient step

        J01.append( self.err(X,Y) ) # evaluate the current error rate

        ## TODO: compute surrogate loss (logistic negative log-likelihood)
        ## Jsurr = sum_i [ (log si) if yi==1 else (log(1-si)) ]
        if (YY[i] == 1):
            Jsurr = np.sum(np.log(sigmoid(ri)))
        else:
            Jsurr = np.sum(np.log(1 - sigmoid(ri)))
        d = Jsurr/M
        Jnll.append(d) # TODO evaluate the current NLL loss
        display.clear_output(wait=True);
        plt.subplot(1,2,1); plt.cla(); plt.plot(Jnll, 'b-', J01, 'r-'); # plot losses
        if N==2: plt.subplot(1,2,2); plt.cla(); self.plotBoundary(X,Y); # & predictor if 2D
        plt.pause(.01); # let OS draw the plot

        ## For debugging: you may want to print current parameters & losses
        # print self.theta, ' => ', Jnll[-1], ' / ', J01[-1]
        # raw_input() # pause for keystroke

        # TODO check stopping criteria: exit if exceeded # of epochs ( > stopEpochs)
        if (epoch > stopEpochs):
            done = True
        if (len(Jnll) > 1):
            if np.abs(Jnll[-2] - Jnll[-1]) < stopTol:
                done = True
            # or if Jnll not changing between epochs ( < stopTol )
```

1.7:

In [26]:

```
class logisticClassify2(ml.classifier):
    classes = []
```

```


theta = np.array( [-1, 0, 0] ) # initialize theta to something
plotBoundary = myPlotBoundary #
predict = myPredict # Now all parts are implemented
train = myTrain

```


```

learnerA = logisticClassify2()
learnerA.theta = np.array([0.,0.,0.]);
learnerA.train(XA,YA,initStep=1e-1,stopEpochs=1000,stopTol=1e-5);
ml.plotClassify2D(learnerA,XA,YA)
print("Training error rate: ",learnerA.err(XA,YA))
plt.show()

```



('Training error rate: ', 0.010101010101010102)




In [27]:


```

learnerB = logisticClassify2()
learnerB.theta = np.array([0.,0.,0.]);
learnerB.train(XB,YB,initStep=1e-1,stopEpochs=1000,stopTol=1e-5);
ml.plotClassify2D(learnerB,XB,YB)
print("Training error rate: ",learnerB.err(XB,YB))
plt.show()

```



('Training error rate: ', 0.25252525252525254)



Problem 2:

1. $T(a + bx_1)$: A and B can be shattered because the points can be separated into +1 and -1. C and D cannot be because there are arrangements for the dots being the same or different that don't allow them to be split into +1 and -1.

2. $T((a \cdot b)x_1 + (c/a)x_2)$: This learner is similar to the first in the sense that it is a line. It can separate A and B but can't separate C and D. A line has to be able to split the data in +1 and -1 and for C and D, points that are similar can end up on the same side which means they can't be shattered.

3. $T((x_1 - a)^2 + (x_2 - b)^2 + c)$: A, B, and C can be shattered because a circle can separate the data points into ones that are similar and different. D can't be shattered because if the last 3 points are the same and the first is different, then a circle won't effectively separate the data.

Problem 3: Statement of Collaboration

I asked a few questions on Piazza and looked at the others' questions and responses on Piazza. I also asked a friend who's good at math to explain gradient descent to me and she did.